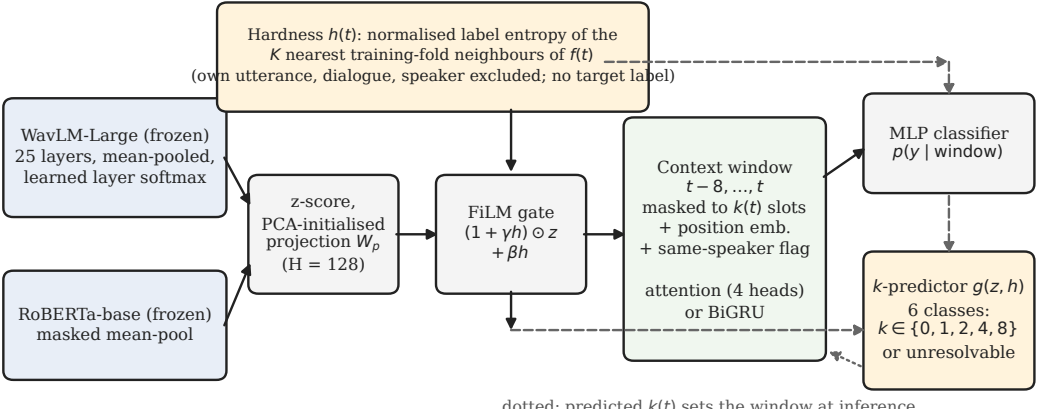
Hardness $h(t)$: normalised label entropy of the K nearest training-fold neighbours of $f(t)$
(own utterance, dialogue, speaker excluded; no target label)
WavLM-Large (frozen)
25 layers, mean-pooled,
learned layer softmax
RoBERTa-base (frozen)
masked mean-pool
z-score,
PCA-initialised
projection W_p
(H = 128)
FiLM gate
$(1+\gamma h)\odot z$
$+\beta h$
Context window
$t-8, \ldots, t$
masked to $k(t)$ slots
+ position emb.
+ same-speaker flag
attention (4 heads)
or BiGRU
MLP classifier
$p(y \mid \text{window})$
k-predictor $g(z, h)$
6 classes:
$k \in \{0, 1, 2, 4, 8\}$
or unresolvable
dotted: predicted $k(t)$ sets the window at inference
dashed: inputs to g; g is supervised by the oracle k^* (inner CV, training folds only)
encoders frozen; 0.32 M trainable parameters below them